



Aalto University
School of Science

Implementing a Virtual Network System among Containers

Songlin Jiang
songlin.jiang@aalto.fi

Tutor: Tuomas Aura

April 21, 2023

Introduction

- Rise of **container technologies** has attracted attention from industry and academia as applications are moving to the cloud.
- **Virtual machines** are commonly used for building and testing network systems configurations before deployed into real world.
- There are many drawbacks in our usage for choosing **virtual machine**, and **container** is a good cure for those drawbacks.
- My paper investigates the possibility of implementing virtual network systems using **Docker containers** instead.



Overview

Introduction

Docker Networking System

- Network Drivers

- Routing, Firewall and IPv6

Case Study: VPN

- VPN Overview

- Host to Host

- Site to Site in IPv6

Evaluation

- Usability and Portability

- Performance

- Security and Limitations

Conclusion

Q&A



Network Drivers

- Docker employs a pluggable networking subsystem.
- Default drivers for Docker networking system include the **bridge**, **host**, **overlay**, **IPVLAN**, **MACVLAN**, and **none**.
- Possible candidates for implementing the virtual network system on one host machine are **bridge**, **IPVLAN**, and **MACVLAN**.
 - The **none** driver disables all networking.
 - The **host** driver shares the same network stack with the host.
 - The **overlay** driver creates a distributed network system among multiple Docker daemon hosts.

MACVLAN

We choose **MACVLAN** finally:

Table: Network Drivers Comparison

Items	bridge	IPVLAN	MACVLAN
Resources	High	Low	Lowest
MAC Address	Different	Same	Different
VM Migration	No	No	Yes

- MACVLAN is a trivial bridge that does not need to learn as it already knows every **MAC address** it can receive.
- MACVLAN also supports **address configuration** and **discovery protocols**, as well as other **multicast protocols**.


Routing, Firewall and IPv6

- By default, Docker does not allow manipulating container network devices and setting routing tables or firewalls inside containers.
 - However, **net_admin** capability in Linux allows us to:
 1. **Make interface configuration.**
 2. **Administrate IP firewall, masquerading, and accounting.**
 3. **Modify routing tables.**
 4. Bind to any address for transparent proxying.
 5. Set the type-of-service (TOS).
 6. Clear driver statistics.
 7. Set the promiscuous mode.
 8. Enable multicasting.
 9. Use setsockopt(2) to set several socket options.
 - **IPv6** is also supported in Docker containers.
-


VPN Overview

- Simulate IoT devices in **sites A, B** connect to server in **cloud S**.
 - **Site A, B**, and **cloud S** both use the private IP addresses to improve security and save the IPv4 addresses.
 - **Gateways A, B, S** connect **sites A, B, S** to the public Internet.
 - The **router** represents the Internet between the sites and cloud.
 - The address space between the **gateway** and **router** simulates public, routable IPv4 addresses, although they are all private.
 - **Site A, B**, and **cloud S** use the **router** to access the Internet.
- We experiment with 2 types of VPN using strongSwan (IPsec):
 1. **Site to Site**
 2. **Host to Host**

Host to Host



Site to Site in IPv6



Usability and Portability

With **Docker Compose**, the Docker orchestration tool, containers **outperforms** VirtualBox-based Vagrant in many aspects:

Table: Functionalities Comparison

Items	Vagrant + VirtualBox	Docker Compose
Resources	Heavy	Lightweight
Kernel	Own	Shared (Namespaces)
Scalability	Hard	Easy
M1/M2 Support	Limited	Fully
Image Hub	Unavailable	Available (Docker Hub)
Seamless	No	Yes



Performance

Table below shows the metrics for the performance evaluation:

Table: Performance Test Result in Average

Solution	Boot Time ¹	Memory ²
Docker Compose	75 s	278 MB
Vagrant + VirtualBox	689 s	4.5 GB

The container based solution *reduces*:

- 1. Fresh boot time¹** by nearly 90%
- 2. Memory consumption²** by nearly 94%

¹Also include the running environment building time for the host platform

²Maximum value during the whole running process

Security and Limitations

- We can't visit the host network interfaces inside containers, and vice versa. The container network stacks are **isolated** from host.
- Limitations of the Docker networking model include:
 1. Disallowing to assign **overlapped IP address ranges** by Docker, even for network interfaces that won't directly connected.
 2. Disallowing to assign **IP address ending in ".1"** by Docker, as these addresses are reserved by Docker for gateways or routers.
- However, we can always choose to configure the IP addresses manually inside containers to bypass these limitations.

Conclusion

- Containers are much more **lightweight** than virtual machines.
- As a mature virtualization technology, containers can **realize every functionality** we require for virtual network systems.
- Container is a suitable **replacement** for virtual machines if you want to migrate the virtual network systems away from VMs.
- Hope it can inspire researchers and engineers to migrate their network systems testing environment from VMs into containers.



Thanks for listening!

Any Questions?



Figure: Scan to get the case study implementation source code
